

Backdoors to Tractable Answer-Set Programming^{*†}

Johannes Klaus Fichte and Stefan Szeider

Vienna University of Technology, Austria

fichte@kr.tuwien.ac.at, stefan@szeider.net

Abstract

We present a unifying approach to the efficient evaluation of propositional answer-set programs. Our approach is based on backdoors which are small sets of atoms that represent “clever reasoning shortcuts” through the search space. The concept of backdoors is widely used in the areas of propositional satisfiability and constraint satisfaction. We show how this concept can be adapted to the nonmonotonic setting and how it allows to augment various known tractable subproblems, such as the evaluation of Horn and acyclic programs.

In order to use backdoors we need to find them first. We utilize recent advances in fixed-parameter algorithmics to detect small backdoors. This implies fixed-parameter tractability of the evaluation of propositional answer-set programs, parameterized by the size of backdoors. Hence backdoor size provides a structural parameter similar to the treewidth parameter previously considered. We show that backdoor size and treewidth are incomparable, hence there are instances that are hard for one and easy for the other parameter. We complement our theoretical results with first empirical results.

1 Introduction

Answer-Set Programming (ASP) is an increasingly popular framework for declarative programming [Marek and Truszczyński, 1999; Niemelä, 1999]. ASP allows to describe a problem by means of rules and constraints that form a disjunctive logic program. Solutions to the program are so-called stable models or answer sets. Many important problems of AI and reasoning can be represented and successfully solved within the ASP framework. However, the main computational problems for ASP (such as deciding whether a program has a solution, or if a certain atom is contained in at least one or in all solutions) are of high worst-case complexity and are located at the second level of the Polynomial Hierarchy [Eiter and Gottlob, 1995]. The known complexity results do not rule out the possibility for exact algorithms that work efficiently for real-world instances by exploiting the presence of a “hidden structure.”

In this paper we follow a new approach of making the vague notion of a hidden structure precise. Our approach is based on the concept of *backdoors* which is widely used in the areas of propositional satisfiability and constraint satisfaction (see, e.g., [Williams *et al.*, 2003a; Gottlob and Szeider, 2006; Samer and Szeider, 2009]), and also for quantified Boolean formulas and argumentation [Samer and Szeider, 2009a; Ordyniak and Szeider, 2011].

A backdoor is a small set of key atoms that represent a “clever reasoning shortcut” through the search space. By deciding the status of the atoms in the backdoor, we can reduce a given program to several tractable programs belonging to a *target class* of programs. Consequently the evaluation of the given program is *fixed-parameter tractable* in the size of the backdoor, i.e., polynomial for fixed backdoor size k where the order of the polynomial is independent of k [Downey and Fellows, 1999]. By allowing backdoors of increasing size $k = 1, 2, 3, \dots$ we can gradually augment a known tractable class of programs.

^{*}Research supported by ERC (COMPLEX REASON 239962).

[†]This is an extended and updated version of a paper that appeared in the Proceedings of IJCAI’11.

Our results are as follows:

- We show that the most important computational problems of propositional answer-set programming, including credulous/skeptical reasoning (and even counting all answer sets) are fixed-parameter tractable in the size of the backdoor.
- We show that the detection of backdoors is fixed-parameter tractable for various target classes, including the class of all Horn programs and classes based on various notions of acyclicity. This way we make recent results of fixed-parameter algorithmics accessible to the field of answer-set programming.
- We show that the concept of backdoors entails fixed-parameter tractability results for answer-set programming [Ben-Eliyahu, 1996] and so provides a unifying framework.
- We compare backdoor size with respect to various base classes with each other and with the recently studied parameter incidence treewidth [Jakl *et al.*, 2009].
- We present first empirical results where we consider the backdoor size of structured programs and random programs of varied density.

2 Formal Background

We consider a universe U of propositional *atoms*. A *literal* is an atom $a \in U$ or its negation $\neg a$. A *disjunctive logic program* (or simply a *program*) P is a set of *rules* of the form $x_1 \vee \dots \vee x_l \leftarrow y_1, \dots, y_n, \neg z_1, \dots, \neg z_m$ where $x_1, \dots, x_l, y_1, \dots, y_n, z_1, \dots, z_m$ are atoms and l, n, m are non-negative integers. We write $\{x_1, \dots, x_l\} = H(r)$ (the *head* of r) and $\{y_1, \dots, y_n, z_1, \dots, z_m\} = B(r)$ (the *body* of r), $B^+(r) = \{y_1, \dots, y_n\}$ and $B^-(r) = \{z_1, \dots, z_m\}$. We denote the sets of atoms occurring in a rule r or in a program P by $\text{at}(r) = H(r) \cup B(r)$ and $\text{at}(P) = \bigcup_{r \in P} \text{at}(r)$, respectively.

A rule r is *negation-free* if $B^-(r) = \emptyset$, r is *normal* if $|H(r)| = 1$, r is a *constraint* if $|H(r)| = 0$, r is *disjunction-free* if $|H(r)| \leq 1$, r is *Horn* if it is negation-free and disjunction-free, and r is *tautological* if $B^+(r) \cap (H(r) \cup B^-(r)) \neq \emptyset$. We say that a program has a certain property if all its rules have the property. We denote by **Horn** the classes of all Horn programs.

A set M of atoms *satisfies* a rule r if $(H(r) \cup B^-(r)) \cap M \neq \emptyset$ or $B^+(r) \setminus M \neq \emptyset$. M is a *model* of P if it satisfies all rules of P . The *GL reduct* of a program P under a set M of atoms is the program P^M obtained from P by first removing all rules r with $B^-(r) \cap M \neq \emptyset$ and second removing all $\neg z$ where $z \in B^-(r)$ from all remaining rules r [Gelfond and Lifschitz, 1991]. M is an *answer set* (or *stable set*) of a program P if M is a minimal model of P^M . We denote by $\text{AS}(P)$ the set of all answer sets of P .

We also need some notions from *propositional satisfiability*. A *clause* is a finite set of literals, a CNF formula is a finite set of clauses. A *truth assignment* is a mapping $\tau : X \rightarrow \{0, 1\}$ defined for a set $X \subseteq U$ of atoms. For $x \in X$ we put $\tau(\neg x) = 1 - \tau(x)$. By $\text{ta}(X)$ we denote the set of all truth assignments $\tau : X \rightarrow \{0, 1\}$. The *truth assignment reduct* of a CNF formula F with respect to $\tau \in \text{ta}(X)$ is the CNF formula F_τ obtained from F by first removing all clauses c that contain a literal set to 1 by τ , and second removing from the remaining clauses all literals set to 0 by τ . τ *satisfies* F if $F_\tau = \emptyset$, and F is *satisfiable* (in symbols $\text{sat}(F)$) if it is satisfied by some τ .

2.1 ASP Problems

The main computational problems for ASP are as follows. **CONSISTENCY**: given a program P , does P have an answer-set? **CREDULOUS/SKEPTICAL REASONING**: given a program P and an atom $a \in \text{at}(P)$, is a contained in some/all answer-set(s) of P ? **AS COUNTING**: how many answer sets does P have? **AS ENUMERATION**: list all answer sets of P . **CONSISTENCY** and **CREDULOUS REASONING** are Σ_2^P -complete, **SKEPTICAL REASONING** is Π_2^P -complete [Eiter and Gottlob, 1995]. The problems remain NP (or co-NP) hard for normal programs [Marek and Truszczyński, 1991], but are polynomial-time solvable for Horn

programs [Gelfond and Lifschitz, 1988]. AS COUNTING is easily seen to be $\#P$ -hard even for normal programs in view of the $\#P$ -completeness of $\#SAT$.

2.2 Fixed-Parameter Tractability

We give some basic background on parameterized complexity. For more detailed information we refer to other sources [Downey and Fellows, 1999; Gottlob and Szeider, 2006]. A *parameterized problem* L is a subset of $\Sigma^* \times \mathbb{N}$ for some finite alphabet Σ . For an instance $(I, k) \in \Sigma^* \times \mathbb{N}$ we call I the *main part* and k the *parameter*. L is *fixed-parameter tractable* if there exist a computable function f and a constant c such that we can decide whether $(I, k) \in L$ in time $O(f(k)\|I\|^c)$ where $\|I\|$ denotes the size of I . FPT is the class of all fixed-parameter tractable decision problems. The *Weft Hierarchy* consists of parameterized complexity classes $W[1] \subseteq W[2] \subseteq \dots$ which are defined as the closure of certain parameterized problems under parameterized reductions. There is strong theoretical evidence that parameterized problems that are hard for classes $W[i]$ are not fixed-parameter tractable.

3 Backdoors

Before we introduce the notion of backdoors to the ASP domain we review it in the domain where it originates from.

3.1 Satisfiability Backdoors

Let F be a CNF formula and X a set of atoms. The following is obvious from the definitions:

(*) F is satisfiable if and only if F_τ is satisfiable for at least one truth assignment $\tau \in \text{ta}(X)$.

This observation leads to the definition of a strong backdoor relative to a class \mathcal{C} of polynomially solvable CNF formulas: a set X of atoms is a *strong \mathcal{C} -backdoor* of a CNF formula F if $F_\tau \in \mathcal{C}$ for all truth assignments $\tau \in \text{ta}(X)$. Assume that the satisfiability of formulas $F \in \mathcal{C}$ of size $\|F\| = n$ can be decided in time $O(n^c)$. Then we can decide the satisfiability of an arbitrary formula F for which we know a strong \mathcal{C} -backdoor of size k in time $O(2^k n^c)$ which is efficient as long as k remains small.

Before we can use the strong backdoor we need to find it first. For most reasonable target classes \mathcal{C} the detection of a strong backdoor of size at most k is NP-hard if k is part of the input. However, as we are interested in finding small backdoors, it makes sense to parameterize the backdoor search by k and consider the parameterized complexity of backdoor detection. Indeed, with respect to the classes of Horn CNF formulas and 2-CNF formulas, the detection of strong backdoors of size $\leq k$ is fixed-parameter tractable [Nishimura *et al.*, 2004; Samer and Szeider, 2009]. For other target classes (clustering formulas and renamable Horn formulas) the detection of deletion backdoors (a subclass of strong backdoors) of size at most k is fixed-parameter tractable [Nishimura *et al.*, 2007; Razgon and O’Sullivan, 2008].

3.2 ASP Backdoors

In order to translate the notion of backdoors to the domain of ASP, we first need to come up with a suitable concept of a reduction with respect to a truth assignment. The following is a natural definition which generalizes a concept of Gottlob *et al.* [2002].

Definition 1. Let P be a program, X a set of atoms, and $\tau \in \text{ta}(X)$. The truth assignment reduct of P under τ is the logic program P_τ obtained from P by

1. removing all rules r with $H(r) \cap \tau^{-1}(1) \neq \emptyset$ or $H(r) \subseteq X$;

2. removing all rules r with $B^+(r) \cap \tau^{-1}(0) \neq \emptyset$;
3. removing all rules r with $B^-(r) \cap \tau^{-1}(1) \neq \emptyset$;
4. removing from the heads and bodies of the remaining rules all literals $v, \neg v$ with $v \in X$.

Definition 2. Let \mathcal{C} be a class of programs. A set X of atoms is a strong \mathcal{C} -backdoor of a program P if $P_\tau \in \mathcal{C}$ for all truth assignments $\tau \in \text{ta}(X)$.

Example 1. Consider the program $P = \{s \leftarrow w; u \leftarrow s, q; r \leftarrow w, s; t \leftarrow \neg r; q \leftarrow \neg s, u; w \leftarrow \neg r, u\}$. The set $X = \{r, s\}$ is a strong **Horn**-backdoor since all four truth assignment reducts $P_{r=0, s=0} = P_{00} = \{t \leftarrow; q \leftarrow u; w \leftarrow u\}$, $P_{01} = \{u \leftarrow q; t \leftarrow; w \leftarrow u\}$, $P_{10} = \{q \leftarrow u\}$, and $P_{11} = \{u \leftarrow q\}$ are Horn programs.

A direct equivalence similar to $(*)$ does not hold for ASP, even if we consider the most basic problem CONSISTENCY. Take for example the program $P = \{x \leftarrow \neg x; y \leftarrow\}$ and the set $X = \{x\}$. Both reducts $P_{x=0} = \{y\}$ and $P_{x=1} = \{y\}$ have answer sets, but P has no answer set. However, we can show a somewhat weaker asymmetric variant of $(*)$, where we can map each answer set of P to an answer set of P_τ for some $\tau \in \text{ta}(X)$. This is made precise by the following definition and lemma.

Definition 3. Let P be a program and X a set of atoms. We define

$$\text{AS}(P, X) = \{M \cup \tau^{-1}(1) : \tau \in \text{ta}(X \cap \text{at}(P)), M \in \text{AS}(P_\tau)\}.$$

Lemma 1. $\text{AS}(P) \subseteq \text{AS}(P, X)$ holds for every program P and every set X of atoms.

Proof. Let $M \in \text{AS}(P)$ be chosen arbitrarily. We put $X_0 = (X \setminus M) \cap \text{at}(P)$ and $X_1 = X \cap M$ and define a truth assignment $\tau \in \text{ta}(X \cap \text{at}(P))$ by setting $\tau^{-1}(i) = X_i$ for $i \in \{0, 1\}$. Let $M' = M \setminus X_1$. Observe that $M' \in \text{AS}(P_\tau)$ implies $M \in \text{AS}(P, X)$ since $M = M' \cup \tau^{-1}(1)$ by definition. Hence, to establish the lemma, it suffices to show that $M' \in \text{AS}(P_\tau)$. We have to show that M' is a model of $P_\tau^{M'}$, and that no proper subset of M' is a model of $P_\tau^{M'}$.

In order to show that M' is a model of $P_\tau^{M'}$, choose $r' \in P_\tau^{M'}$ arbitrarily. By construction of $P_\tau^{M'}$ there is a corresponding rule $r \in P$ with $H(r') = H(r) \setminus X_0$ and $B^+(r') = B^+(r) \setminus X_1$ which gives rise to a rule $r'' \in P_\tau$, and in turn, r'' gives rise to $r' \in P_\tau^{M'}$. Since $B^-(r) \cap X_1 = \emptyset$ (otherwise r would have been deleted forming P_τ) and $B^-(r) \cap M' = \emptyset$ (otherwise r'' would have been deleted forming $P_\tau^{M'}$), it follows that $B^-(r) \cap M = \emptyset$. Thus r gives rise to a rule $r^* \in P^M$ with $H(r) = H(r^*)$ and $B^+(r) = B^+(r^*)$. Since $M \in \text{AS}(P)$, M satisfies r^* , i.e., $H(r) \cap M \neq \emptyset$ or $B^+(r) \setminus M \neq \emptyset$. However, $H(r) \cap M = H(r') \cap M'$ and $B^+(r) \setminus M = B^+(r') \setminus M'$, thus M' satisfies r' . Since $r' \in P_\tau^{M'}$ was chosen arbitrarily, we conclude that M' is a model of $P_\tau^{M'}$.

In order to show that no proper subset of M' is a model of $P_\tau^{M'}$ choose arbitrarily a proper subset $N' \subsetneq M'$. Let $N = N' \cup X_1$. Since $M' = M \setminus X_1$ and $X_1 \subseteq M$ it follows that $N \subsetneq M$. Since M is a minimal model of P^M , N cannot be a model of P^M . Consequently, there must be a rule $r \in P$ such that $B^-(r) \cap M = \emptyset$ (i.e., r is not deleted by forming P^M), $B^+(r) \subseteq N$ and $H(r) \cap N = \emptyset$. However, since M satisfies P^M , and since $B^+(r) \subseteq N \subseteq M$, $H(r) \cap M \neq \emptyset$. Thus r is not a constraint. Moreover, since $H(r) \cap M \neq \emptyset$ and $M \cap X_0 = \emptyset$, it follows that $H(r) \setminus X_0 \neq \emptyset$. Thus, since $H(r) \cap X_1 = \emptyset$, $H(r) \setminus X \neq \emptyset$. We conclude that r is not deleted when forming P_τ and giving rise to a rule $r' \in P_\tau$, which in turn is not deleted when forming $P_\tau^{M'}$, giving rise to a rule r'' , with $H(r'') = H(r) \setminus X_0$, $B^+(r'') = B^+(r) \setminus X_1$, and $B^-(r'') = \emptyset$. Since $B^+(r'') \subseteq N'$ and $H(r'') \cap N = \emptyset$, N' is not a model of $P_\tau^{M'}$.

Thus we have established that M' is a stable model of P_τ , and so the lemma follows. \square

In view of Lemmas 1 and 2, we can compute $\text{AS}(P)$ by (i) computing $\text{AS}(P_\tau)$ for all $\tau \in \text{ta}(X)$ (this produces the set $\text{AS}(P, X)$ of candidates for $\text{AS}(P)$), and (ii) checking for each $M \in \text{AS}(P, X)$ whether it is an answer-set of P . The check (ii) entails (iia) checking whether $M \in \text{AS}(P, X)$ is a model of P and (iib) whether $M \in \text{AS}(P, X)$ is a minimal model of P^M . We would like to note that in particular

any constraint contained in P is removed in the truth assignment reduct P_τ but considered in check (iia). Clearly check (iia) can be carried out in polynomial time for each M . Check (iib), however is co-NP-hard in general [Eiter and Gottlob, 1995], but polynomial for normal programs [Cadoli and Lenzerini, 1994]. Fortunately, for our considerations it suffices to perform check (iib) for programs that are “close to Horn”, and so the check is fixed-parameter tractable in the size of the given backdoor, as we shall show in the following lemma¹.

Lemma 2. *Let \mathcal{C} be a class of normal programs. Given a program P , a strong \mathcal{C} -backdoor set X of P , and $M \in \text{AS}(P, X)$. Then deciding whether M is an answer set of P is fixed-parameter tractable for parameter $|X|$. In particular, this decision can be made in time $O(2^k n)$ where n denotes the input size of P and $k = |X|$.*

Proof. Let \mathcal{C} be a class of normal programs, P a program, and X a strong \mathcal{C} -backdoor set X of P with $|X| = k$. We can check in polynomial time whether M is a model of P^M . If it is not, we can reject M , and we are done. Hence assume that M is a model of P^M . In order to check whether $M \in \text{AS}(P)$ we still need to decide whether M is a minimal model of P^M . We may assume, w.l.o.g., that P contains no tautological rules, as it is clear that the test for minimality does not depend on tautological rules.

Let $X_1 \subseteq M \cap X$. We construct from P^M a program $P_{X_1 \subseteq X}^M$ by (i) removing all rules r for which $H(r) \cap X_1 \neq \emptyset$, and (ii) replacing for all remaining rules r the head $H(r)$ with $H(r) \setminus X$, and the positive body $B^+(r)$ with $B^+(r) \setminus X_1$.

Claim: $P_{X_1 \subseteq X}^M$ is Horn.

To show the claim, consider some rule $r' \in P_{X_1 \subseteq X}^M$. By construction, there must be a rule $r \in P$ that gives rise to a rule in P^M , which in turn gives rise to r' . Let $\tau \in \text{ta}(X)$ be the assignment that sets all atoms in $X \cap H(r)$ to 0, and all atoms in $X \setminus H(r)$ to 1. Since r is not tautological, it follows that r is not deleted when we obtain P_τ , and it gives rise to a rule $r^* \in P_\tau$, where $H(r^*) = H(r) \setminus X$. However, since \mathcal{C} is a class of normal programs, r^* is normal. Hence $1 \geq |H(r^*)| = |H(r) \setminus X| = |H(r')|$, and the claim follows.

To test whether M is a minimal model of P^M , we run the following procedure for every set $X_1 \subseteq M \cap X$.

If $P_{X_1 \subseteq X}^M$ has no model, then stop and return TRUE. Otherwise, compute the unique minimal model L of the Horn program $P_{X_1 \subseteq X}^M$. If $L \subseteq M \setminus X$, $L \cup X_1 \subsetneq M$, and $L \cup X_1$ is a model of P^M , then return FALSE. Otherwise return TRUE.

For each set $X_1 \subseteq M \cap X$ the above procedure runs in linear time. This follows directly from the fact that we can compute the unique minimal model of a Horn program in linear time [Dowling and Gallier, 1984]. As there are $O(2^k)$ sets X_1 to consider, we have a total running time of $O(2^k n)$ where n denotes the input size of P and $k = |X|$. It remains to establish the correctness of the algorithm in terms of the following claim.

Claim: M is a minimal model of P^M if and only if the algorithm returns TRUE for each $X_1 \subseteq M \cap X$.

(\Rightarrow). Assume that M is a minimal model of P^M , and suppose to the contrary that there is some $X_1 \subseteq M \cap X$ for which the algorithm returns FALSE. Consequently, $P_{X_1 \subseteq X}^M$ has a unique minimal model L with $L \subseteq M \setminus X$, $L \cup X_1 \subsetneq M$, and where $L \cup X_1$ is a model of P^M . This contradicts the assumption that M is a minimal model of P^M . Hence the only-if direction of the lemma is shown.

(\Leftarrow). Assume that the algorithm returns TRUE for each $X_1 \subseteq M \cap X$. We show that M is a minimal model of P^M . Suppose to the contrary that P^M has a model $M' \subsetneq M$.

We run the algorithm for $X_1 := M' \cap X$. By assumption, the algorithm returns TRUE. There are two possibilities: (i) $P_{X_1 \subseteq X}^M$ has no model, or (ii) $P_{X_1 \subseteq X}^M$ has a model, and for its unique minimal model L the following holds: L is not a subset of $M \setminus X$, or $L \cup X_1$ is not a proper subset of M , or $L \cup X_1$ is not a model of P^M .

We show that case (i) is not possible, by showing that $M' \setminus X$ is a model of $P_{X_1 \subseteq X}^M$.

¹In [Fichte and Szeider, 2011] we overlooked that the minimality check requires some additional attention.

To see this, consider a rule $r' \in P_{X_1 \subseteq X}^M$, and let $r \in P^M$ such that r' is obtained from r by removing X from $H(r)$ and by removing X_1 from $B^+(r)$. Since M' is a model of P^M , we have (a) $B^+(r) \setminus M' \neq \emptyset$ or (b) $H(r) \cap M' \neq \emptyset$. Moreover, since $B^+(r') = B^+(r) \setminus X_1$ and $X_1 = M' \cap X$, (i) implies $\emptyset \neq B^+(r) \setminus M' = B^+(r) \setminus X_1 \setminus M' = B^+(r') \setminus M' \subseteq B^+(r') \setminus (M' \setminus X)$, and since $H(r) \cap X_1 = \emptyset$, (ii) implies $\emptyset \neq H(r) \cap M' = H(r) \cap (M' \setminus X_1) = H(r) \cap (M' \setminus X) = (H(r) \setminus X) \cap (M' \setminus X) = H(r') \cap (M' \setminus X)$. Hence $M' \setminus X$ satisfies r' . Since $r' \in P_{X_1 \subseteq X}^M$ was chosen arbitrarily, we conclude that $M' \setminus X$ is a model of $P_{X_1 \subseteq X}^M$.

Case (ii) is not possible either, as we can see as follows. Assume $P_{X_1 \subseteq X}^M$ has a model, and let L be its unique minimal model. Since $M' \setminus X$ is a model of $P_{X_1 \subseteq X}^M$, as shown above, we have $L \subseteq M' \setminus X$.

We have $L \subseteq M \setminus X$ since $L \subseteq M' \setminus X$ and $M' \setminus X \subseteq M \setminus X$.

Further we have $L \cup X_1 \subsetneq M$ since $L \cup X_1 \subseteq (M' \setminus X) \cup X_1 = (M' \setminus X) \cup (M' \cap X) = M' \subsetneq M$.

And finally $L \cup X_1$ is a model of P^M , as can be seen as follows. Consider a rule $r \in P^M$. If $X_1 \cap H(r) \neq \emptyset$, then $L \cup X_1$ satisfies r ; thus it remains to consider the case $X_1 \cap H(r) = \emptyset$. In this case there is a rule $r' \in P_{X_1 \subseteq X}^M$ with $H(r') = H(r) \setminus X$ and $B^+(r') = B^+(r) \setminus X_1$. Since L is a model of $P_{X_1 \subseteq X}^M$, L satisfies r' . Hence (a) $B^+(r') \setminus L \neq \emptyset$ or (b) $H(r') \cap L \neq \emptyset$. Since $B^+(r') = B^+(r) \setminus X_1$, (a) implies that $B^+(r) \setminus (L \cup X_1) \neq \emptyset$; and since $H(r') \subseteq H(r)$, (b) implies that $H(r) \cap (L \cup X_1) \neq \emptyset$. Thus $L \cup X_1$ satisfies r . Since $r \in P^M$ was chosen arbitrarily, we conclude that $L \cup X_1$ is a model of P^M .

Since neither case (i) nor case (ii) is possible, we have a contradiction, and we conclude that M is a minimal model of P^M .

Hence the second direction of the claim is established, and so the lemma follows. \square

Thus, in view of Lemmas 1 and 2, the computation of $\text{AS}(P)$ is fixed-parameter tractable for parameter k if we know a strong \mathcal{C} -backdoor set X of size at most k for P , and each program in \mathcal{C} is normal and its stable sets can be computed in polynomial time. This consideration leads to the following definition and result.

Definition 4. A class \mathcal{C} of programs is enumerable if for each $P \in \mathcal{C}$ we can compute $\text{AS}(P)$ in polynomial time.

For any class \mathcal{C} of programs we denote by \mathcal{C}^* the class containing all programs that belong to \mathcal{C} after removal of tautological rules and constraints. It is easy to see that whenever \mathcal{C} is enumerable, then so is \mathcal{C}^* . Note that all classes considered in this paper are enumerable.

Theorem 1. Let \mathcal{C} be an enumerable class of normal programs. Problems CONSISTENCY, CREDULOUS and SKEPTICAL REASONING, AS COUNTING and AS ENUMERATION are all fixed-parameter tractable when parameterized by the size of a strong \mathcal{C} -backdoor, assuming that the backdoor is given as an input.

Proof. Let X be the given backdoor, $k = |X|$ and n the input size of P . Since $P_\tau \in \mathcal{C}$ and \mathcal{C} is enumerable, we can compute $\text{AS}(P_\tau)$ in polynomial time for each $\tau \in \text{ta}(X)$, say in time $O(n^c)$. Observe that therefore $|\text{AS}(P_\tau)| \leq O(n^c)$ for each $\tau \in \text{ta}(X)$. Thus we obtain $\text{AS}(P, X)$ in time $O(2^k n^c)$, and $|\text{AS}(P, X)| \leq O(2^k n^c)$. By Lemma 1, $\text{AS}(P) \subseteq \text{AS}(P, X)$. By means of Lemma 2 we can decide whether $M \in \text{AS}(P)$ in time $O(2^k n)$ for each $M \in \text{AS}(P, X)$. Thus we determine from $\text{AS}(P, X)$ the set of all answer sets of P in time $O(2^k \cdot n^c \cdot 2^k \cdot n + 2^k \cdot n^c) = O(2^{2k} n^{c+1})$. Once we know $\text{AS}(P)$, then we can also answer any of the listed problems within polynomial time. \square

If we know that each program in \mathcal{C} has at most one answer set, and P has a strong \mathcal{C} -backdoor of size k , then we can conclude that P has at most 2^k answer sets. Thus, we obtain an upper bound on the number of answer sets of P by computing a small strong \mathcal{C} -backdoor of P .

Example 2. We consider program P of Example 1. The answer sets of P_τ are $\text{AS}(P_{00}) = \{\{t\}\}$, $\text{AS}(P_{01}) = \{\{t\}\}$, $\text{AS}(P_{10}) = \{\emptyset\}$, and $\text{AS}(P_{11}) = \{\emptyset\}$ for $\tau \in \text{ta}(\{r, s\})$. $\text{AS}(P, X) = \{\{t\}, \{t, s\}, \{r\}, \{r, s\}\}$, and since only $\{t\} \in \text{AS}(P, X)$ is an answer set of P , we obtain $\text{AS}(P) = \{\{t\}\}$.

3.3 Deletion Backdoors

We will see below that the following variant of strong backdoors is often useful. For a program P and a set X of atoms we define $P - X$ as the program obtained from P by deleting all atoms contained in X from all the rules (heads and bodies) of P .

Definition 5. *Let \mathcal{C} be a class of programs. A set X of atoms is a deletion \mathcal{C} -backdoor of a program P if $P - X \in \mathcal{C}$.*

In general, not every strong \mathcal{C} -backdoor is a deletion \mathcal{C} -backdoor, and not every deletion \mathcal{C} -backdoor is a strong \mathcal{C} -backdoor. We call \mathcal{C} to be *rule induced* if for each $P \in \mathcal{C}$, $P' \subseteq P$ implies $P' \in \mathcal{C}$. Note that many natural classes of programs (and all classes considered in this paper) are rule induced.

Proposition 1. *If \mathcal{C} is rule induced, then every deletion \mathcal{C} -backdoor is a strong \mathcal{C} -backdoor.*

Proof. The statement follows from the fact that $P_\tau \subseteq P - X$ for every $\tau \in \text{ta}(X)$ and every program P . \square

3.4 Backdoor Detection

In order to use Theorem 1 we need to find the backdoor first. Each class \mathcal{C} of programs gives rise to the following parameterized problem: **STRONG \mathcal{C} -BACKDOOR DETECTION**: given a program P and an integer k , find a strong \mathcal{C} -backdoor X of P of size at most k , or report that such X does not exist. We also consider the problem **DELETION \mathcal{C} -BACKDOOR DETECTION**, defined similarly (which is in some cases easier to solve).

4 Target Class Horn

We first consider the important case **Horn*** as the target class for backdoors. It is well known that normal Horn programs have a unique answer set and this set can be found in linear time [Dowling and Gallier, 1984], hence **Horn** and **Horn*** are enumerable. The following lemma shows that **Horn*** is particularly well suited as a target class.

Lemma 3. *A set X is a strong **Horn***-backdoor of a program P if and only if it is a deletion **Horn***-backdoor of P .*

Proof. It suffices to show the lemma for **Horn**, and in view of Proposition 1 it suffices to show the only-if direction. Assume X is a strong **Horn**-backdoor of P . Consider a rule $r' \in P - X$ which is neither tautological nor a constraint. Let $r \in P$ be a rule from which r' was obtained in forming $P - X$. We define $\tau \in \text{ta}(X)$ by setting all atoms in $H(r) \cup B^-(r)$ to 0, all atoms in $B^+(r)$ to 1, and all remaining atoms in $X \setminus \text{at}(r)$ arbitrarily to 0 or 1. Since r is not tautological, this definition of τ is sound. It remains to observe that $r' \in P_\tau$. Since X is a strong **Horn**-backdoor of P , the rule r' is Horn. \square

Theorem 2. **STRONG **Horn***-BACKDOOR DETECTION** is fixed-parameter tractable in k . In fact, given a program P with n atoms we can find a strong **Horn***-backdoor of size $\leq k$ in time $O(1.2738^k + kn)$ or decide that no such backdoor exists.

Proof. (Sketch) Let G be the undirected graph defined on the set of atoms of the given program P , where two atoms x, y are joined by an edge if and only if P contains a non-tautological rule r with $x, y \in H(r)$ or $x \in H(r)$ and $y \in B^-(r)$. Now it is easy to see that a set $X \subseteq \text{at}(P)$ is a vertex cover of G if and only if X is a deletion **Horn***-backdoor of P . A vertex cover of size $\leq k$, if it exists, can be found in time $O(1.2738^k + kn)$ [Chen *et al.*, 2006]. The theorem follows by Lemma 3. \square

For instance, the undirected graph G of the program P of Example 1 consists of the two paths (w, r, t) and (s, q) . Then $\{r, s\}$ is a vertex cover of G . We observe easily that there exists no vertex cover of size 1. Thus $\{r, s\}$ is a smallest strong **Horn***-backdoor of P .

Ben-Eliyahu [1996] showed that evaluation of normal logic programs is fixed-parameter tractable in (i) the number of atoms that appear in negative rule bodies, and (ii) the total number of non-Horn rules. It is not difficult to see that both numbers are greater or equal to the size of a smallest strong **Horn***-backdoor, and so entailed by our approach.

5 Acyclicity-Based Target Classes

There are two causes for a program to have a large number of answer sets: (i) disjunctions in the heads of rules, and (ii) certain cyclic dependencies between rules. Disallowing both causes yields so-called *stratified* programs [Gelfond and Lifschitz, 1988]. In the following we will study backdoor detection for various classes of stratified programs. We define the classes by requiring normality and acyclicity (the absence of certain types cycles). In order to define acyclicity we associate with each program P its *directed dependency graph* D_P and its *undirected dependency graph* U_P where D_P is an extended version of the dependency graph in [Apt *et al.*, 1988] and U_P of the undirected dependency graph in [Gottlob *et al.*, 2002]. D_P has as vertices the atoms of P , a directed edge (x, y) between any two atoms x, y for which there is a rule $r \in P$ with $x \in H(r)$ and $y \in B(r)$ or a rule $r \in P$ with $x, y \in H(r)$; if there is a rule $r \in P$ with $x \in H(r)$ and $y \in B^-(r)$ or there is a rule $r \in P$ with $x, y \in H(r)$, then the edge (x, y) is called a *negative edge*. U_P is obtained from D_P by replacing each negative edge $e = (x, y)$ with two undirected edges $\{x, v_e\}, \{v_e, y\}$ where v_e is a new *negative vertex*, and by replacing each remaining directed edge (u, v) with an undirected edge $\{u, v\}$. By a *directed cycle* of P we mean a directed cycle in D_P , by an *undirected cycle* of P we mean an undirected cycle in U_P . Figure 1 visualizes D_P and U_P of the program P of Example 1. A directed (undirected) cycle is *bad* if it contains a negative edge (a negative vertex), otherwise it is *good*. Various classes of programs arise by requiring the programs to have no directed bad cycles (**DBC-Acyc**), no undirected bad cycles (**BC-Acyc**), no directed cycles (**DC-Acyc**), and no undirected cycles (**C-Acyc**). **DBC-Acyc**, the largest class among the considered classes, is exactly the class **Strat** of stratified programs [Apt *et al.*, 1988].

For instance in the program P of Example 1, $(r, v_{(w,r)}, w, r)$ is an undirected cycle, (u, q, u) is a directed cycle, $(s, v_{(q,s)}, q, u, s)$ is an undirected bad cycle, and (w, r, w) is a directed bad cycle (see Figure 1).




Figure 1: Directed dependency graph D_P (left) and undirected dependency graph U_P (right) of the program P of Example 1.

In order to compare the size of backdoors for the various classes, we need to compare the classes themselves (evidently, if $\mathcal{C} \subsetneq \mathcal{C}'$, then every strong (deletion) \mathcal{C}' -backdoor is also a strong (deletion) \mathcal{C} -backdoor, but not necessarily the other way around). By definition we have **DC-Acyc** \subsetneq **DBC-Acyc** and **C-Acyc** \subsetneq **BC-Acyc** \subsetneq **DBC-Acyc**; it is easy to see that the inclusions are proper. However, contrary to what one expects, **C-Acyc** $\not\subseteq$ **DC-Acyc**, which can be seen by considering the program $P_0 = \{x \leftarrow y, y \leftarrow x\}$, hence **C-Acyc** and **DC-Acyc** are incomparable. Requiring that a program has no directed cycles but may have directed good cycles of length 2 (as in P_0) gives rise to the class **DC2-Acyc**,

which generalizes both classes **C-Acyc** and **DC-Acyc**. The diagram in Figure 2 shows the relationship between the various program classes.




Figure 2: Relationship between classes of programs. An arrow from \mathcal{C} to \mathcal{C}' indicates that \mathcal{C} is a proper subset of \mathcal{C}' . If there is no arrow between two classes (or the arrow does not follow by transitivity of set inclusion), then the two classes are incomparable.

Theorem 3. *For each class $\mathcal{C} \in \{\mathbf{C-Acyc}, \mathbf{BC-Acyc}, \mathbf{DC-Acyc}, \mathbf{DC2-Acyc}, \mathbf{Strat}\}$ the problem STRONG \mathcal{C} -BACKDOOR DETECTION is $W[2]$ -hard and therefore unlikely to be fixed-parameter tractable.*

Proof. (Sketch) We give a reduction from the $W[2]$ -complete problem HITTING SET [Downey and Fellows, 1999]. An instance of this problem is a pair (S, k) where $S = \{S_1, \dots, S_m\}$ is a family of sets and k is an integer (the parameter). The question is whether there exists a set H of size at most k which intersects with all the S_i ; such H is a hitting set. We construct a program P as follows. As atoms we take the elements of $X = \bigcup_{i=1}^m S_i$ and new atoms a_i^j and b_i^j for $1 \leq i \leq m$, $1 \leq j \leq k+1$. For each $1 \leq i \leq m$ and $1 \leq j \leq k+1$ we take two rules r_i^j, s_i^j where $H(r_i^j) = \{a_i^j\}$, $B^-(r_i^j) = S_i$, $B^+(r_i^j) = S_i \cup \{b_i^j\}$; $H(s_i^j) = \{b_i^j\}$, $B^-(s_i^j) = \{a_i^j\}$, $B^+(s_i^j) = \emptyset$. The result now follows by showing that S has a hitting set of size $\leq k$ if and only if P has a strong \mathcal{C} -backdoor of size $\leq k$ where \mathcal{C} is any of the classes mentioned. \square

For **DC-Acyc**, **DC2-Acyc**, and **Strat** we can avoid the use of tautological rules in the reduction and so strengthen Theorem 3 as follows (it would be interesting to know if this is also possible for the remaining two classes mentioned in Theorem 3).

Theorem 4. *For each class $\mathcal{C} \in \{\mathbf{DC-Acyc}, \mathbf{DC2-Acyc}, \mathbf{Strat}\}$ the problem STRONG \mathcal{C}^* -BACKDOOR DETECTION is $W[2]$ -hard and therefore unlikely to be fixed-parameter tractable.*

Proof. (Sketch) We modify the above reduction from HITTING SET by redefining the rules r_i^j, s_i^j . We put $H(r_i^j) = \{a_i^j\}$, $B^-(r_i^j) = S_i$, $B^+(r_i^j) = \{b_i^j\}$; $H(s_i^j) = \{b_i^j\}$, $B^-(s_i^j) = \{a_i^j\}$, $B^+(s_i^j) = X$. \square

The $W[2]$ -hardness results suggests to relax the considered problems and look for *deletion* backdoors: Which of the classes mentioned in Theorem 3 admit fixed-parameter tractable detection of deletion backdoors? Using very recent results from fixed-parameter algorithmics we can answer this question positively for all considered classes except for **Strat** whose complexity remains open.

The $W[2]$ -hardness results suggest to relax the considered problems and to look for *deletion* backdoors: Which of the classes mentioned in Theorem 3 admit fixed-parameter tractable detection of deletion backdoors? Using very recent results from fixed-parameter algorithmics we can answer this question positively for all considered classes except for **Strat** whose complexity remains open.

Theorem 5. *For each class $\mathcal{C} \in \{\mathbf{C-Acyc}, \mathbf{BC-Acyc}, \mathbf{DC-Acyc}, \mathbf{DC2-Acyc}\}$ the problem DELETION \mathcal{C}^* -BACKDOOR DETECTION is fixed-parameter tractable.*

Proof. (1) DELETION **C-Acyc**-BACKDOOR DETECTION can be solved by solving the FEEDBACK VERTEX SET (FVS) problem for U_p , which is well-known to be FPT [Downey and Fellows, 1999], this was already observed by Gottlob *et al.* [2002]. (2) DELETION **DC-Acyc**-BACKDOOR DETECTION is equivalent to the DIRECTED FVS problem on D_p . The parameterized complexity of DIRECTED FVS remained open for many years and was recently shown FPT by Chen *et al.* [2008] with a break-through result. (3) DELETION

BC-Acyc-BACKDOOR DETECTION can be solved by solving the EDGE SUBSET FVS problem, where only cycles need to be covered that contain an edge from a given set S . This problem was recently shown FPT by Cygan *et al.* [2010] and Kawarabayashi and Kobayashi [2010]. (4) Finally, the problem DELETION **DC2-Acyc**-BACKDOOR DETECTION can be solved by finding a feedback vertex set in a *mixed graph* (a graph containing directed and undirected edges), as we can replace a good cycle on two edges by one undirected edge. FVS FOR MIXED GRAPHS was recently shown FPT by Bonsma and Lokshtanov [2010]. \square

The classes mentioned in Theorem 3 are rule-induced. Hence we can use this theorem to strengthen the fixed-parameter tractability result of Theorem 1 by dropping the assumption that the backdoor is given.

The considered classes can be generalized by taking the parity of the number of negative edges on bad cycles into account. In recent research Fichte [2011] generalized the tractability results of Lin and Zhao [2004] by considering backdoors with respect to such parity classes.

6 Theoretical Comparison of Parameters

In this section we compare ASP parameters in terms of their generality.

Jakl *et al.* [2009] applied the graph parameter treewidth to ASP and showed that the main reasoning problems for ASP are fixed-parameter tractable by the treewidth of the *incidence graph* of the program. The incidence graph of a program P is the bipartite graph on the rules and atoms of P , where a rule and an atom are joined by an edge if and only if the atom occurs in the rule. It turns out that incidence treewidth is incomparable with backdoor size for various target classes considered.

Theorem 6. *Let $\mathcal{C} \in \{\mathbf{Horn}, \mathbf{C-Acyc}, \mathbf{DC-Acyc}, \mathbf{DC2-Acyc}, \mathbf{BC-Acyc}, \mathbf{Strat}\}$. There are programs whose incidence graphs have constant treewidth but require arbitrarily large strong and deletion \mathcal{C} -backdoors, and there are programs where the converse prevails.*

Proof. (Sketch) Clearly each of the considered classes contains programs of arbitrary high treewidth of the incidence graph. Conversely, we consider a program $P \notin \mathcal{C}$. We denote by nP the program consisting of the union of n atom-disjoint copies of P . By basic properties of treewidth it follows that the treewidth of the incidence graph of P equals to that of nP , however, smallest \mathcal{C} -backdoors are of size $\geq n$. \square

The treewidth approach is based on dynamic programming which is of high space complexity and therefore only practical for instances of treewidth below 10 [Jakl *et al.*, 2009]. The backdoor approach is more space efficient since for each partial truth assignment $\tau \in \text{ta}(X)$ of a backdoor X , the computations of $\text{AS}(P_\tau)$ and the corresponding elements of $\text{AS}(P_\tau)$ can be carried out independently.

One might ask whether it makes sense to consider restrictions on the treewidth of the undirected dependency graph, defined above. However, this restriction does not yield tractability, as the reduction of [Eiter and Gottlob, 1995] produces programs with undirected dependency graphs of treewidth 2.

7 Empirical Comparison of Parameters

We have determined the size of smallest backdoors for various programs, including *structured* real-world instances and *random* instances. The results are summarized in Table 1. It is known that so-called tight programs are closely related to SAT [Lin and Zhao, 2003]. The QueensEqTest instances and two of the Daimler-Chrysler instances are tight, all other instances considered are not tight. For pragmatic reasons we have used **Horn*** as the target class as smallest backdoors are easy to compute, even for large inputs. Our experimental results indicate that structured instances have smaller backdoors than random instances. It also seems that random instances with higher density have larger backdoors.

We have conducted a second series of experiments on random instances where we have analyzed how much we gain by considering the more general acyclicity-based target classes instead of **Horn**. It appears that smallest deletion **Horn***-backdoors are indeed significantly larger than deletion backdoors for

instance set	vars	bd (%)	stdev
Daimler-Chrysler-MT	1785.14	21.46	1.56
Daimler-Chrysler-NC	1793.0	22.94	2.92
Daimler-Chrysler-RZ	1562.5	11.53	2.21
Daimler-Chrysler-SZ	1567.53	13.97	3.47
Daimler-Chrysler-UC	1781.74	21.4	2.01
Daimler-Chrysler-UT	1781.23	23.53	3.95
Mutex	6449.0	49.94	0.09
RLP-3	150.0	58.28	1.37
RLP-4	150.0	64.53	0.92
RLP-5	150.0	68.4	0.97
RLP-6	150.0	70.9	0.82
RLP-7	150.0	73.68	0.87
RLP-8	150.0	75.54	0.74
RG-40	40.0	93.5	1.24
RG-50	50.0	94.05	0.96
RG-60	60.0	94.38	0.82

Table 1: Size of smallest strong **Horn**-backdoors (bd) for various benchmark sets, given as % of the total number of variables (vars) by the mean over the instances. Daimler-Chrysler-⟨test⟩: 554 Real-world instances encoding logistics problems from car configurations. The disjunctive programs have been compiled from SAT instances provided by Sinz et al. [2003] grouped by the kind of consistency test. The instances are produced using the simple encoding where a clause $\{a, b, \neg c, \neg d\}$ becomes the rule $a, b \leftarrow c, d$. Mutex: Disjunctive programs that encode the equivalence test of partial implementations of circuits, provided by Maratea et al. [2008] based on QBF instances of Ayari and Basin [2000] RLP-⟨ ρ ⟩: Randomly generated normal programs provided by Zhao and Lin [2003] of various density ρ (number of rules divided by the number of variables) with 10 instances per step. RG-⟨ n ⟩: Randomly generated instances provided by Gebser [Asp, 2009] with $n = 40, 50$, and 60 variables, respectively with 40 instances per step.

acyclicity-based target classes. The distinction between directed and undirected cycles seems to have a significant effect on the backdoor size, whereas the distinction between good and bad cycles seems to be less significant. However these results are not fully conclusive as the considered programs were rather small.

8 Conclusion

We have introduced the backdoor approach to the domain of propositional answer-set programming. The backdoor approach allows to augment a known tractable class and makes the efficient solving methods for the tractable class generally applicable. Our approach makes recent results in fixed-parameter algorithmics applicable to nonmonotonic reasoning. The comparison results show that the parameters based on backdoor size are incomparable with treewidth and therefore provide fixed-parameter tractability for programs that are hard for the treewidth approach.

The results and concepts of this paper give rise to several interesting research questions. For instance, it would be interesting to consider backdoors for target classes that contain programs with an exponential number of answer-sets, but where the set of all answer-sets can be succinctly represented. A simple example is the class of programs that consist of independent components of bounded size. Other questions are concerned with alternative ways of using backdoors. For instance, by means of “backdoor trees” [Samer

and Szeider, 2008] one can avoid the consideration of all 2^k partial assignments of the backdoor and thus make the backdoor approach feasible for programs with larger backdoors. A further use of backdoors that seems worth exploring is the control of heuristics of ASP solvers.

References

- [Apt *et al.*, 1988] K. R. Apt, H. A. Blair, and A. Walker. *Towards a theory of declarative knowledge*, Morgan Kaufmann, 1988.
- [Asp, 2009] Asparagus. <http://asparagus.cs.uni-potsdam.de>, 2009.
- [Ayari and Basin, 2000] A. Ayari and D. Basin. Bounded model construction for monadic second-order logics. *CAV 2000*.
- [Ben-Eliyahu, 1996] R. Ben-Eliyahu. A hierarchy of tractable subsets for computing stable models. *J. Artif. Intell. Res.*, 5, 1996.
- [Bonsma and Lokshtanov, 2010] P. Bonsma and D. Lokshtanov. Feedback vertex set in mixed graphs. *Algorithms and Data Structures*, 122–133, Springer, 2011.
- [Cadoli and Lenzerini, 1994] M. Cadoli and M. Lenzerini. The complexity of propositional closed world reasoning and circumscription. *Journal of Computer and System Sciences*, 48(2), 255–310, Elsevier, 1994.
- [Chen *et al.*, 2006] J. Chen, I. A. Kanj, and G. Xia. Improved parameterized upper bounds for vertex cover. *MFCS 2006*.
- [Chen *et al.*, 2008] J. Chen, Y. Liu, S. Lu, B. O’Sullivan, and I. Razgon. A fixed-parameter algorithm for the directed feedback vertex set problem. *J. ACM*, 55(5), 2008.
- [Cygan *et al.*, 2010] M. Cygan, M. Pilipczuk, M. Pilipczuk, and J. O. Wojtaszczyk. Subset feedback vertex set is fixed parameter tractable. arXiv 1004.2972, 2010.
- [Dowling and Gallier, 1984] W. F. Dowling and J. H. Gallier. Linear-time algorithms for testing the satisfiability of propositional Horn formulae. *J. Logic Programming*, 1(3), 1984.
- [Downey and Fellows, 1999] R. G. Downey and M. R. Fellows. *Parameterized Complexity*. Springer, 1999.
- [Eiter and Gottlob, 1995] T. Eiter and G. Gottlob. On the computational cost of disjunctive logic programming: propositional case. *Ann. Math. Artif. Intell.*, 15(3-4), 1995.
- [Fichte and Szeider, 2011] Fichte, J.K., Szeider, S.: Backdoors to tractable answer-set programming. *IJCAI 2011*.
- [Fichte, 2011] J. K. Fichte. The good, the bad, the odd – cycles in answer set programming. *ESSLLI 2011 Student Session*, 2011.
- [Gelfond and Lifschitz, 1988] M. Gelfond and V. Lifschitz. The stable model semantics for logic programming. *ICLP 1988*.
- [Gelfond and Lifschitz, 1991] M. Gelfond and V. Lifschitz. Classical negation in logic programs and disjunctive databases. *New Generation Comput.*, 9(3/4), 1991.
- [Gottlob and Szeider, 2006] G. Gottlob and S. Szeider. Fixed-parameter algorithms for artificial intelligence, constraint satisfaction, and database problems. *The Computer Journal*, 51(3), 2006.

- [Gottlob *et al.*, 2002] G. Gottlob, F. Scarcello, and M. Sideri. Fixed-parameter complexity in AI and nonmonotonic reasoning. *Artificial Intelligence*, 138(1-2), 2002.
- [Jakl *et al.*, 2009] M. Jakl, R. Pichler, and S. Woltran. Answer-set programming with bounded treewidth. *IJCAI 2009*.
- [Janhunen and Oikarinen, 2002] T. Janhunen and E. Oikarinen. Testing the equivalence of logic programs under stable model semantics. *Logics in AI*, 2002.
- [Kawarabayashi and Kobayashi, 2010] K. Kawarabayashi and Y. Kobayashi. Fixed-parameter tractability for the subset feedback set problem and the s-cycle packing problem. Technical report, University of Tokyo, Japan, 2010.
- [Lin and Zhao, 2003] F. Lin and J. Zhao. On tight logic programs and yet another translation from normal logic programs to propositional logic. *IJCAI 2003*.
- [Lin and Zhao, 2004] F. Lin and X. Zhao. On odd and even cycles in normal logic programs. *AAAI 2004*.
- [Maratea *et al.*, 2008] M. Maratea, F. Ricca, W. Faber, and N. Leone. Look-back techniques and heuristics in DLV: Implementation, evaluation, and comparison to QBF solvers. *J. Algorithms*, 63, 2008.
- [Marek and Truszczyński, 1991] V. W. Marek and M. Truszczyński. Autoepistemic logic. *J. ACM*, 38(3), 1991.
- [Marek and Truszczyński, 1999] V. W. Marek and M. Truszczyński. Stable models and an alternative logic programming paradigm. *The Logic Programming Paradigm: a 25-Year Perspective*, 1999.
- [Niemelä, 1999] I. Niemelä. Logic programs with stable model semantics as a constraint programming paradigm. *Ann. Math. Artif. Intell.*, 25(3-4), 1999.
- [Nishimura *et al.*, 2004] N. Nishimura, P. Ragde, and S. Szeider. Detecting backdoor sets with respect to Horn and binary clauses. *SAT 2004*.
- [Nishimura *et al.*, 2007] N. Nishimura, P. Ragde, and S. Szeider. Solving #SAT using vertex covers. *Acta Informatica*, 44(7-8), 2007.
- [Ordyniak and Szeider, 2011] S. Ordyniak and S. Szeider. Augmenting Tractable Fragments of Abstract Argumentation. *IJCAI 2011*.
- [Razgon and O’Sullivan, 2008] I. Razgon and B. O’Sullivan. Almost 2-sat is fixed-parameter tractable (extended abstract). *ICALP 2008*.
- [Samer and Szeider, 2008] M. Samer and S. Szeider. Backdoor trees. *AAAI 2008*.
- [Samer and Szeider, 2009] M. Samer and S. Szeider. Fixed-parameter tractability. *Handbook of Satisfiability*, ch. 13, 2009.
- [Samer and Szeider, 2009a] M. Samer and S. Szeider. Backdoor sets of quantified Boolean formulas. *JAR*, 42(1):77–97, 2009.
- [Sinz *et al.*, 2003] C. Sinz, A. Kaiser, and W. Küchlin. Formal methods for the validation of automotive product configuration data. *AI EDAM 2003*.
- [Williams *et al.*, 2003a] R. Williams, C. Gomes, and B. Selman. Backdoors to typical case complexity. *IJCAI 2003*.
- [Zhao and Lin, 2003] Y. Zhao and F. Lin. Answer set programming phase transition: A study on randomly generated programs. *ICLP 2003*.